

Algorithm Design for Big Data Systems

Divide and conquer

Graph algorithms

Streaming algorithms



Embarrassingly Parallel Problems

- Problems that can be easily decomposed into many independent problems.
- Examples.
 - Word count
 - k-means
 - PageRank

Divide and Conquer

code at:

<https://www.cse.ust.hk/msbd5003/nb/DC.ipynb>



Classical Divide-and-Conquer

- Classical D&C
 - Divide problem into 2 parts
 - Recursively solve each part
 - Combine the results together
- D&C under big data systems
 - Divide problem into p partitions, where (ideally) p is the number of executors in the system
 - Solve the problem on each partition
 - Combine the results together
- Example: `sum()`, `reduce()`

Prefix Sums

- Input: Sequence x of n elements, binary associative operator $+$
- Output: Sequence y of n elements, with
$$y_k = x_1 + \dots + x_k$$
- Example:
$$x = [1, 4, 3, 5, 6, 7, 0, 1]$$
$$y = [1, 5, 8, 13, 19, 26, 26, 27]$$
- Algorithm:
 - Compute sum for each partition
 - Compute the prefix sums of the p sums
 - Compute prefix sums in each partition

Variants of Prefix Sums

- Assign consecutive id's for each element
 - `zipWithIndex()`
- Given a list of words, find the first appearance of “spark”
- Given two long strings, compare them lexicographically
- Given a sequence of integers, check whether these numbers are monotonically decreasing.

Sorting (Sample Sort)

- Step 1: Sampling
 - Master node collects a sample of sp elements (will determine s later)
- Step 2: Choose splitters
 - Pick every $(i \cdot s)$ -th element in the sample as splitters, $i = 1, \dots, p - 1$
 - Broadcast them to all machines
- Step 3: Shuffling
 - Each machine partitions its data using the splitters
 - Send data to the target machine
- Step 4: Sort each partition
 - Each machine sorts all data received

Probability Tools

- Chernoff inequality

Theorem: Let X_1, X_2, \dots, X_n be independent 0-1 random variables (not necessarily identical), and let $X = \sum_i X_i$ with $\mu = E[X]$. Then for any $0 \leq \delta \leq 1$,

$$\Pr[X \leq (1 - \delta)\mu] \leq \exp\left(-\frac{\mu\delta^2}{2}\right)$$

$$\Pr[X \geq (1 + \delta)\mu] \leq \exp\left(-\frac{\mu\delta^2}{3}\right)$$

- Union bound. Given events E_1, \dots, E_n , independent or not,

$$\Pr\left[\bigcup_{i=1}^n E_i\right] \leq \sum_{i=1}^n \Pr[E_i]$$

Determining Sample Size

- Goal: No machine receives more than $(1 + \epsilon) \frac{N}{p}$ elements w.h.p.
 - How large should s be?
- Let the elements be a_1, \dots, a_N in sorted order
- A sub-sequence $a_i, \dots, a_{i+(1+\epsilon)\frac{N}{p}-1}$ is **bad** if it contains $< s$ sampled elements
 - Goal achieved if no sub-sequence is bad
- Consider a particular sub-sequence
 - $X = \#$ sampled elements in it; $E[X] = \frac{sp}{N} \cdot (1 + \epsilon) \frac{N}{p} = (1 + \epsilon)s$
 - By Chernoff inequality:
$$\Pr[X < s] \leq \Pr\left[X < \left(1 - \frac{\epsilon}{2}\right) E[X]\right] \leq e^{-\epsilon^2 s / 8}$$
- By union bound, $\Pr[\exists \text{ a bad subsequence}] \leq N \cdot e^{-\epsilon^2 s / 8}$
 - If want failure probability 1%, suffices to set $s = \frac{8 \ln(100N)}{\epsilon^2}$
 - A tighter analysis improves the $\log N$ term to $\log \frac{p}{\epsilon}$.

Distributed Sampling

- Q: How to sample one element uniformly from n elements stored on p servers?
- A:
 - First randomly sample a server
 - Then ask that server to return an element randomly chosen from its n/p elements.
 - The probability of each element being sampled is $\frac{1}{p} \cdot \frac{p}{n} = \frac{1}{n}$
- Q: How to sample many elements at once?
- A: Do each of the two steps above in batch mode
 - First sample sp servers with replacement (this can be done at the master node).
 - If a server is sampled k times, we ask that server to return k samples (with replacement) from its local data.

The Maximum Subarray Problem

Input: Profit history of a company of the years.

Year	1	2	3	4	5	6	7	8	9
Profit (M\$)	-3	2	1	-4	5	2	-1	3	-1

Problem: Find the span of years in which the company earned the most

Answer: Year 5-8 , 9 M\$

Formal definition:

Input: An array of numbers $A[1 \dots n]$, both positive and negative

Output: Find the maximum $V(i, j)$, where $V(i, j) = \sum_{k=i}^j A[k]$

A divide-and-conquer algorithm

Year	1	2	3	4	5	6	7	8	9
Profit (M\$)	-3	2	1	-4	5	2	-1	3	-1

Idea:

Cut the array into two halves

All subarrays can be classified into three cases:

- Case 1: entirely in the first half
- Case 2: entirely in the second half
- Case 3: crosses the cut

Largest of three cases is final solution

The optimal solutions for case 1 and 2 can be found recursively.

Only need to consider case 3.

Solving case 3

Year	1	2	3	4	5	6	7	8	9
Profit (M\$)	-3	2	1	-4	5	2	-1	3	-1



Idea:

Let $q = \lfloor (p + r)/2 \rfloor$

Any case 3 subarray must have starting position $\leq q$, and ending position $\geq q + 1$

Such a subarray can be divided into two parts $A[i..q]$ and $A[q + 1..j]$, for some i and j

Just need to maximize each of them separately

Maximize $A[i..q]$ and $A[q + 1, j]$:

Let i' , j' , be the indices that maximize the values.

i' , j' can be found using linear scans to left and right of q

$A[i', j']$ has largest value of all subarrays that cross q

The (binary) divide-and-conquer algorithm

MaxSubarray (A, p, r) :

```
if  $p = r$  then return  $A[p]$ 
 $q \leftarrow \lfloor (p + r) / 2 \rfloor$ 
 $M_1 \leftarrow \text{MaxSubarray}(A, p, q)$ 
 $M_2 \leftarrow \text{MaxSubarray}(A, q + 1, r)$ 
 $L_m \leftarrow -\infty, R_m \leftarrow -\infty$ 
 $V \leftarrow 0$ 
for  $i \leftarrow q$  downto  $p$ 
     $V \leftarrow V + A[i]$ 
    if  $V > L_m$  then  $L_m \leftarrow V$ 
 $V \leftarrow 0$ 
for  $i \leftarrow q + 1$  to  $r$ 
     $V \leftarrow V + A[i]$ 
    if  $V > R_m$  then  $R_m \leftarrow V$ 
return  $\max\{M_1, M_2, L_m + R_m\}$ 
```

First call: **MaxSubarray** ($A, 1, n$)


- Analysis:
 - Recurrence:
$$T(n) = 2T(n/2) + n$$
 - So, $T(n) = \Theta(n \log n)$

If we use the same algorithm on Spark:

- Level 1:
 - Naively: 2 executors are working, all others idle
 - time = $O(n/2)$
 - Smarter: L_m and R_m can be found by the prefix-sum algorithm
 - Can use all executors, time = $O(n/p)$
- Level 2:
 - We have 4 subarrays, and solve two prefix-sums for each subarray
 - Each subarray has size $n/4$, and we make sure that each has the same number of partitions
 - Time = $O(n/p)$
- Level 3: Time = $O(n/p)$
- Stop recursion when each subarray is one partition.
- Total time: $O\left(\frac{n}{p} \cdot \log p\right)$

A linear-time algorithm?

- Define: $X[i] = A[1] + \dots + A[i - 1]$



Year	1	2	3	4	5	6	7	8	9
Profit (M\$)	-3	2	1	-4	5	2	-1	3	-1
$X[i]$		-3	-1	0	-4	1	3	2	5

Observations:


$$V(i, j - 1) = \sum_{k=i}^{j-1} A[k] = X[j] - X[i]$$

For fixed j , finding largest $V(i, j - 1)$ is same as finding the index $i, i < j$ for which $X[i]$ is smallest

Idea: doing this for each j , then find overall largest $V(i, j)$

A linear-time algorithm?

- Define: $X[i] = A[1] + \dots + A[i - 1]$
- Goal: Find $\max_{i < j} (X[j] - X[i])$



Year	1	2	3	4	5	6	7	8	9
Profit (M\$)	-3	2	1	-4	5	2	-1	3	-1
$X[i]$		-3	-1	0	-4	1	3	2	5

Algorithm:

For each j , needs to know $i < j$ that minimizes $X[i]$ (i.e., maximizes $X[j] - X[i]$)

– (Then maximize over all j)

Algorithm increases j by +1 each step

Keeps track of smallest $X[i]$ so far

– Could be old smallest one or it could be current $X[j]$

The linear-time algorithm

```
 $V_{max} \leftarrow -\infty, X_{min} = 0$   
 $X \leftarrow 0, V \leftarrow 0$   
for  $i \leftarrow 1$  to  $n$  do  
     $V \leftarrow V + A[i]$   
    if  $V > V_{max}$  then  $V_{max} \leftarrow V$   
     $X \leftarrow X + A[i]$   
    if  $X < X_{min}$  then  
         $X_{min} \leftarrow X$   
         $V \leftarrow 0$   
return  $V_{max}$ 
```

X_{min} keeps track of smallest $X[i]$ so far.

V contains difference between current $X[i]$ and smallest $X[i]$ so far

Even “simpler”:

```
 $V_{max} \leftarrow -\infty, V \leftarrow 0$   
for  $i \leftarrow 1$  to  $n$  do  
     $V \leftarrow V + A[i]$   
    if  $V > V_{max}$  then  $V_{max} \leftarrow V$   
    if  $V < 0$  then  $V \leftarrow 0$   
return  $V_{max}$ 
```

- Observation:
 - $X < X_{min}$ iff $V < 0$
 - Because $V = X - X_{min}$
 - No need to actually store X !

A more efficient algorithm

- For each partition, solve the problem directly using one executor and the linear-time algorithm
- Now it remains to solve the “cross the boundary” case
- Find the L_m and R_m for each partition, as well as its sum
- For each contiguous subsets of partitions $(i, j), i < j$, the optimal solution with left boundary in partition i and right boundary in partition j is
$$L_m[i] + Sum[i + 1, \dots, j - 1] + R_m[j]$$
- Total time: $O(n/p)$